

Lecture Pre-read materials

A good understanding of attention mechanisms and transformers is expected for this lecture. Please go through the following resources in advance.

Essential Pre-reads-

- Jurafsky & Martin — SLP Chapter 13 — RNNs, vanishing gradients, LSTMs, Seq2Seq, and Attention: web.stanford.edu
- Illustrated Transformer — A highly recommended blog for building intuition around transformer architecture: jalammar.github.io
- Attention Mechanisms — A comprehensive blog covering the foundations of attention: lilianweng.github.io
- Transformers and Pre-training — A video lecture covering key concepts: [Watch on YouTube](#)
- Short Notes — A concise reference document for quick revision: [View Notes](#)

Additional Pre-reads (*for those who wish to explore further*)

For a deeper, research-level understanding, participants may refer to the foundational papers on:

- Cross Attention — [Bahdanau et al., 2014](#)
- Self Attention — [Vaswani et al., 2017](#)